

# COE59410- Generative Deep Learning

## Homework 1

Imran A. Zualkernan

The purpose of this homework is to learn how to use the GPU machines to run large models. In addition, build familiarity with the Keras framework and the associated tools.

## Deliverables

**You need to upload the following on ilearn (2 items)**

1. The Jupyter notebook in its original format.
2. A PDF of the Jupyter notebook for grading.

*Please do not upload a zipped file.* Upload each file separately. Each question is worth 25 points.

- Q1. Load and run the large\_scale\_processing v1.1 Jupyter notebook on the GPU machine and show how you can use tensorboard to monitor the runs remotely on your local machine.
- Q2. Modify the model in large\_scale\_processing v1.1 so that rather than a CNN, the model is a fully connected feedforward neural network. Fine tune the model to show your best results. Report and discuss all the results that are necessary to determine the goodness of your best model.

**Hint: Use the Reshape Layer in Keras.**

- Q3. Use the following two call-backs on your best fully connected model and determine if you are able to improve the results. Clearly explain why or why not.
  - i. LearningRateScheduler
  - ii. ReduceLROnPlateau
- Q4. Use the Keras Hypertune and Random optimizers (<https://keras-team.github.io/keras-tuner/>) to determine if you can improve the model by varying the number of layers, neurons in each layer and the learning rate.

- i. Plot the precision vs. recall of the best 20 models in one figure.
- ii. Show a complete evaluation of the top two models.

## Group 2

- Eman ,
- Huangjin Zhou, b00080932
- Mueez ,

```
1 # Useful links
  # https://www.hostinger.com/tutorials/ssh/basic-ssh-commands

2 from tensorflow.python.client import device_lib
  print(device_lib.list_local_devices())

  [name: "/device:CPU:0"
   device_type: "CPU"
   memory_limit: 268435456
   locality {
   }
   incarnation: 11557652930811666046
  , name: "/device:GPU:0"
   device_type: "GPU"
   memory_limit: 7491581376
   locality {
     bus_id: 1
     links {
     }
   }
  ]
  incarnation: 5153159438992320913
  physical_device_desc: "device: 0, name: Quadro RTX 4000, pci bus id: 0000:01:00.0, compute capab.
  ]

3 import tensorflow as tf
  # print(tf.config.list_physical_devices('GPU'))
  tf.config.experimental.list_physical_devices('GPU')
  tf.config.experimental.list_physical_devices(device_type=None)
  tf.test.is_gpu_available()
  print(tf.test.is_built_with_cuda())

  WARNING:tensorflow:From <ipython-input-3-78f31cea5abc>:5: is_gpu_available (from tensorflow.pytho
  Instructions for updating:
  Use `tf.config.list_physical_devices('GPU')` instead.
  True

4 import numpy as np
  import pandas as pd
  import os
  import matplotlib.pyplot as plt
  from IPython.display import Image, display
  import random
  import math
  import keras
  from keras.preprocessing.text import Tokenizer
  from keras.models import Model, Sequential
```

```

from keras.utils import plot_model
from keras.layers import Reshape, Input, Dense, Dropout, Flatten, Activation, Concatenate
from keras.layers import Conv2D, MaxPooling2D, AveragePooling2D
from keras.optimizers import Adam
from keras import backend, models
#import tensorflow_addons as tfa
import tensorflow as tf
print(tf.__version__)

```

```

# need to add these for the GPU
config = tf.compat.v1.ConfigProto()
config.gpu_options.allow_growth = True
session = tf.compat.v1.Session(config=config)

```

2.4.1

```

5 # import the image generator
from tensorflow.keras.preprocessing.image import ImageDataGenerator

```

```

6 #Setting the parameters for training

```

```

# batch size and image width to use
batch_size=128
width=100

```

```

# all the data directories
train_dir='train/'
test_dir='test/'
valid_dir='valid/'

```

```

# the number of epochs
num_epochs=10

```

```

# creating an image generator that will feed the data from
# each of the directories

```

```

# we use scaling transformation in this generator
generator=ImageDataGenerator(rescale=1./255)

```

```

# we specify the size of the input and batch size
# size of the input is necessary because the image
# needs to be rescaled for the neural network

```

```

train_data=generator.flow_from_directory(train_dir, target_size=(width,width),batch_size=batch_size)
valid_data=generator.flow_from_directory(valid_dir, target_size=(width,width),batch_size=batch_size)
test_data=generator.flow_from_directory(test_dir, target_size=(width,width),batch_size=batch_size)

```

```

# the number of steps per epoch is samples/batch size
# we need to use these numbers later

```

```

train_steps_per_epoch=math.ceil(train_data.samples/batch_size)
valid_steps_per_epoch=math.ceil(valid_data.samples/batch_size)
test_steps_per_epoch=math.ceil(test_data.samples/batch_size)
print(train_steps_per_epoch)
print(valid_steps_per_epoch)
print(test_steps_per_epoch)

```

```

Found 35215 images belonging to 250 classes.
Found 1250 images belonging to 250 classes.
Found 1250 images belonging to 250 classes.

```

```

276
10
10

```

```

7 # the actual model should go here
hidden_units = 256*2*2*2*2
dropout = 0.1
num_labels = train_data.num_classes

```

```

model = Sequential()
model.add(Reshape((-1,), input_shape=(width, width, 3)))
model.add(Dense(hidden_units, activation='relu'))
model.add(Dropout(dropout))
model.add(Dense(hidden_units/2, activation='relu'))
model.add(Dropout(dropout))
model.add(Dense(hidden_units/4, activation='relu'))
model.add(Dropout(dropout))
model.add(Dense(num_labels, activation='softmax'))
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
reshape (Reshape)	(None, 30000)	0
dense (Dense)	(None, 8192)	245768192
dropout (Dropout)	(None, 8192)	0
dense_1 (Dense)	(None, 4096)	33558528
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 2048)	8390656
dropout_2 (Dropout)	(None, 2048)	0
dense_3 (Dense)	(None, 250)	512250
=====		
Total params: 288,229,626		
Trainable params: 288,229,626		
Non-trainable params: 0		

```

8 # Compile the model
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

```

```

9 # see if the model is good.
print(model)

```

<tensorflow.python.keras.engine.sequential.Sequential object at 0x7fd3183dc890>

```

10 from tensorflow.keras.callbacks import LearningRateScheduler, ReduceLROnPlateau

```

```

def lr_schedule(epoch):
    """Learning rate scheduler - called every epoch"""
    lr = 1e-3
    fold = int(epoch / 10) + 1
    lr /= fold

    return lr

```

```
lr_scheduler = LearningRateScheduler(lr_schedule)

lr_reducer = ReduceLROnPlateau(factor=np.sqrt(0.1),
                               cooldown=0,
                               patience=5,
                               min_lr=0.5e-6)
```

Q1. Load and run the large\_scale\_processing v1.1 Jupyter notebook on the GPU machine and show how you can use tensorboard to monitor the runs remotely on your local machine.

```
$ jupyter notebook --port 9999 --NotebookApp.allow_remote_access=True
```

```
[I 20:42:04.633 NotebookApp] Serving notebooks from local directory: /home/group2/
[I 20:42:04.633 NotebookApp] Jupyter Notebook 6.2.0 is running at:
[I 20:42:04.633 NotebookApp] http://localhost:9999/?token=7c9d27bf2c24e9e66eee1da6
[I 20:42:04.633 NotebookApp] or http://127.0.0.1:9999/?token=7c9d27bf2c24e9e66eee
[I 20:42:04.633 NotebookApp] Use Control-C to stop this server and shut down all k
[W 20:42:04.635 NotebookApp] No web browser found: could not locate runnable brows
[C 20:42:04.635 NotebookApp]
```

To access the notebook, open this file in a browser:

```
file:///home/group2/.local/share/jupyter/runtime/nbserver-28383-open.html
```

Or copy and paste one of these URLs:

```
http://localhost:9999/?token=7c9d27bf2c24e9e66eee1da6f5b51d4bd2b8669393e49
```

```
or http://127.0.0.1:9999/?token=7c9d27bf2c24e9e66eee1da6f5b51d4bd2b8669393e49
```

```
$ tensorboard --logdir logs/fit --port=8888
```

Serving TensorBoard on localhost; to expose to the network, use a proxy or pass --  
TensorBoard 2.4.0 at http://localhost:8888/ (Press CTRL+C to quit)

```
11 from tensorflow.keras.callbacks import TensorBoard
    tensorboard = TensorBoard(log_dir='logs/fit')

    print(valid_steps_per_epoch)
    num_epochs = 20

    callbacks = [lr_reducer, lr_scheduler, tensorboard]
    history=model.fit(train_data,
                      steps_per_epoch=train_steps_per_epoch,
                      validation_data=valid_data,
                      epochs=num_epochs,
                      validation_steps=valid_steps_per_epoch, callbacks=callbacks)

10
Epoch 1/100
276/276 [=====] - 28s 100ms/step - loss: 17.9556 - accuracy: 0.0076 - v
Epoch 2/100
```

276/276 [=====] - 28s 100ms/step - loss: 5.0856 - accuracy: 0.0239 - val  
Epoch 3/100  
276/276 [=====] - 27s 99ms/step - loss: 4.8207 - accuracy: 0.0412 - val  
Epoch 4/100  
276/276 [=====] - 28s 100ms/step - loss: 4.6110 - accuracy: 0.0584 - val  
Epoch 5/100  
276/276 [=====] - 27s 99ms/step - loss: 4.4302 - accuracy: 0.0761 - val  
Epoch 6/100  
276/276 [=====] - 28s 100ms/step - loss: 4.2868 - accuracy: 0.0929 - val  
Epoch 7/100  
276/276 [=====] - 27s 99ms/step - loss: 4.1800 - accuracy: 0.1043 - val  
Epoch 8/100  
276/276 [=====] - 28s 100ms/step - loss: 4.1198 - accuracy: 0.1146 - val  
Epoch 9/100  
276/276 [=====] - 28s 100ms/step - loss: 4.0230 - accuracy: 0.1205 - val  
Epoch 10/100  
276/276 [=====] - 28s 100ms/step - loss: 3.9548 - accuracy: 0.1318 - val  
Epoch 11/100  
276/276 [=====] - 28s 99ms/step - loss: 3.7698 - accuracy: 0.1620 - val  
Epoch 12/100  
276/276 [=====] - 28s 100ms/step - loss: 3.7020 - accuracy: 0.1675 - val  
Epoch 13/100  
276/276 [=====] - 28s 100ms/step - loss: 3.6439 - accuracy: 0.1771 - val  
Epoch 14/100  
276/276 [=====] - 28s 100ms/step - loss: 3.6025 - accuracy: 0.1799 - val  
Epoch 15/100  
276/276 [=====] - 27s 99ms/step - loss: 3.5496 - accuracy: 0.1902 - val  
Epoch 16/100  
276/276 [=====] - 28s 100ms/step - loss: 3.5025 - accuracy: 0.2007 - val  
Epoch 17/100  
276/276 [=====] - 27s 99ms/step - loss: 3.4508 - accuracy: 0.2055 - val  
Epoch 18/100  
276/276 [=====] - 27s 99ms/step - loss: 3.4350 - accuracy: 0.2077 - val  
Epoch 19/100  
276/276 [=====] - 28s 100ms/step - loss: 3.3848 - accuracy: 0.2167 - val  
Epoch 20/100  
276/276 [=====] - 28s 100ms/step - loss: 3.3431 - accuracy: 0.2215 - val  
Epoch 21/100  
276/276 [=====] - 28s 99ms/step - loss: 3.2387 - accuracy: 0.2357 - val  
Epoch 22/100  
276/276 [=====] - 28s 100ms/step - loss: 3.1623 - accuracy: 0.2508 - val  
Epoch 23/100  
276/276 [=====] - 28s 101ms/step - loss: 3.1490 - accuracy: 0.2579 - val  
Epoch 24/100  
276/276 [=====] - 28s 100ms/step - loss: 3.0807 - accuracy: 0.2633 - val  
Epoch 25/100  
276/276 [=====] - 28s 100ms/step - loss: 3.0757 - accuracy: 0.2644 - val  
Epoch 26/100  
276/276 [=====] - 27s 99ms/step - loss: 3.0393 - accuracy: 0.2738 - val  
Epoch 27/100  
276/276 [=====] - 27s 99ms/step - loss: 3.0046 - accuracy: 0.2759 - val  
Epoch 28/100  
276/276 [=====] - 28s 101ms/step - loss: 2.9574 - accuracy: 0.2868 - val  
Epoch 29/100  
276/276 [=====] - 28s 101ms/step - loss: 2.9143 - accuracy: 0.2944 - val  
Epoch 30/100  
276/276 [=====] - 28s 100ms/step - loss: 2.8867 - accuracy: 0.3022 - val  
Epoch 31/100  
276/276 [=====] - 27s 99ms/step - loss: 2.8026 - accuracy: 0.3138 - val  
Epoch 32/100  
276/276 [=====] - 28s 101ms/step - loss: 2.7576 - accuracy: 0.3221 - val  
Epoch 33/100  
276/276 [=====] - 28s 100ms/step - loss: 2.7267 - accuracy: 0.3302 - val  
Epoch 34/100  
276/276 [=====] - 28s 100ms/step - loss: 2.7007 - accuracy: 0.3363 - val

Epoch 35/100  
276/276 [=====] - 28s 100ms/step - loss: 2.6681 - accuracy: 0.3413 - val  
Epoch 36/100  
276/276 [=====] - 28s 100ms/step - loss: 2.6279 - accuracy: 0.3470 - val  
Epoch 37/100  
276/276 [=====] - 28s 100ms/step - loss: 2.6259 - accuracy: 0.3531 - val  
Epoch 38/100  
276/276 [=====] - 27s 99ms/step - loss: 2.5915 - accuracy: 0.3535 - val  
Epoch 39/100  
276/276 [=====] - 28s 100ms/step - loss: 2.5432 - accuracy: 0.3658 - val  
Epoch 40/100  
276/276 [=====] - 28s 100ms/step - loss: 2.4871 - accuracy: 0.3771 - val  
Epoch 41/100  
276/276 [=====] - 28s 100ms/step - loss: 2.4600 - accuracy: 0.3798 - val  
Epoch 42/100  
276/276 [=====] - 28s 100ms/step - loss: 2.4297 - accuracy: 0.3875 - val  
Epoch 43/100  
276/276 [=====] - 28s 100ms/step - loss: 2.4035 - accuracy: 0.4010 - val  
Epoch 44/100  
276/276 [=====] - 28s 100ms/step - loss: 2.3760 - accuracy: 0.4014 - val  
Epoch 45/100  
276/276 [=====] - 28s 100ms/step - loss: 2.3580 - accuracy: 0.4062 - val  
Epoch 46/100  
276/276 [=====] - 28s 100ms/step - loss: 2.3046 - accuracy: 0.4170 - val  
Epoch 47/100  
276/276 [=====] - 28s 100ms/step - loss: 2.2868 - accuracy: 0.4199 - val  
Epoch 48/100  
276/276 [=====] - 28s 100ms/step - loss: 2.2572 - accuracy: 0.4276 - val  
Epoch 49/100  
276/276 [=====] - 28s 100ms/step - loss: 2.2403 - accuracy: 0.4357 - val  
Epoch 50/100  
276/276 [=====] - 28s 100ms/step - loss: 2.2445 - accuracy: 0.4334 - val  
Epoch 51/100  
276/276 [=====] - 27s 99ms/step - loss: 2.1690 - accuracy: 0.4480 - val  
Epoch 52/100  
276/276 [=====] - 28s 100ms/step - loss: 2.1456 - accuracy: 0.4507 - val  
Epoch 53/100  
276/276 [=====] - 28s 100ms/step - loss: 2.1320 - accuracy: 0.4572 - val  
Epoch 54/100  
276/276 [=====] - 28s 100ms/step - loss: 2.1133 - accuracy: 0.4655 - val  
Epoch 55/100  
276/276 [=====] - 28s 100ms/step - loss: 2.0881 - accuracy: 0.4685 - val  
Epoch 56/100  
276/276 [=====] - 28s 100ms/step - loss: 2.0831 - accuracy: 0.4695 - val  
Epoch 57/100  
276/276 [=====] - 28s 100ms/step - loss: 2.0491 - accuracy: 0.4786 - val  
Epoch 58/100  
276/276 [=====] - 28s 100ms/step - loss: 2.0318 - accuracy: 0.4788 - val  
Epoch 59/100  
276/276 [=====] - 28s 100ms/step - loss: 2.0020 - accuracy: 0.4866 - val  
Epoch 60/100  
276/276 [=====] - 28s 99ms/step - loss: 1.9888 - accuracy: 0.4939 - val  
Epoch 61/100  
276/276 [=====] - 28s 100ms/step - loss: 1.9852 - accuracy: 0.4909 - val  
Epoch 62/100  
276/276 [=====] - 28s 100ms/step - loss: 1.9580 - accuracy: 0.5003 - val  
Epoch 63/100  
276/276 [=====] - 28s 100ms/step - loss: 1.9306 - accuracy: 0.5053 - val  
Epoch 64/100  
276/276 [=====] - 28s 100ms/step - loss: 1.9132 - accuracy: 0.5073 - val  
Epoch 65/100  
276/276 [=====] - 28s 100ms/step - loss: 1.8943 - accuracy: 0.5151 - val  
Epoch 66/100  
276/276 [=====] - 28s 99ms/step - loss: 1.8958 - accuracy: 0.5164 - val  
Epoch 67/100

276/276 [=====] - 28s 100ms/step - loss: 1.8562 - accuracy: 0.5239 - val  
Epoch 68/100  
276/276 [=====] - 28s 100ms/step - loss: 1.8596 - accuracy: 0.5245 - val  
Epoch 69/100  
276/276 [=====] - 27s 99ms/step - loss: 1.8423 - accuracy: 0.5284 - val  
Epoch 70/100  
276/276 [=====] - 28s 100ms/step - loss: 1.8089 - accuracy: 0.5361 - val  
Epoch 71/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7696 - accuracy: 0.5419 - val  
Epoch 72/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7929 - accuracy: 0.5411 - val  
Epoch 73/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7757 - accuracy: 0.5498 - val  
Epoch 74/100  
276/276 [=====] - 27s 99ms/step - loss: 1.7553 - accuracy: 0.5510 - val  
Epoch 75/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7105 - accuracy: 0.5555 - val  
Epoch 76/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7375 - accuracy: 0.5497 - val  
Epoch 77/100  
276/276 [=====] - 28s 100ms/step - loss: 1.7055 - accuracy: 0.5609 - val  
Epoch 78/100  
276/276 [=====] - 28s 99ms/step - loss: 1.6559 - accuracy: 0.5726 - val  
Epoch 79/100  
276/276 [=====] - 28s 100ms/step - loss: 1.6777 - accuracy: 0.5704 - val  
Epoch 80/100  
276/276 [=====] - 28s 101ms/step - loss: 1.6612 - accuracy: 0.5706 - val  
Epoch 81/100  
276/276 [=====] - 28s 100ms/step - loss: 1.6391 - accuracy: 0.5783 - val  
Epoch 82/100  
276/276 [=====] - 28s 100ms/step - loss: 1.6321 - accuracy: 0.5789 - val  
Epoch 83/100  
276/276 [=====] - 28s 100ms/step - loss: 1.6149 - accuracy: 0.5870 - val  
Epoch 84/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5950 - accuracy: 0.5865 - val  
Epoch 85/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5674 - accuracy: 0.5956 - val  
Epoch 86/100  
276/276 [=====] - 27s 99ms/step - loss: 1.5804 - accuracy: 0.5909 - val  
Epoch 87/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5799 - accuracy: 0.5927 - val  
Epoch 88/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5608 - accuracy: 0.5998 - val  
Epoch 89/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5501 - accuracy: 0.6031 - val  
Epoch 90/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5289 - accuracy: 0.6047 - val  
Epoch 91/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5120 - accuracy: 0.6124 - val  
Epoch 92/100  
276/276 [=====] - 28s 100ms/step - loss: 1.5022 - accuracy: 0.6164 - val  
Epoch 93/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4978 - accuracy: 0.6143 - val  
Epoch 94/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4972 - accuracy: 0.6140 - val  
Epoch 95/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4752 - accuracy: 0.6153 - val  
Epoch 96/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4629 - accuracy: 0.6232 - val  
Epoch 97/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4579 - accuracy: 0.6228 - val  
Epoch 98/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4504 - accuracy: 0.6279 - val  
Epoch 99/100  
276/276 [=====] - 28s 100ms/step - loss: 1.4177 - accuracy: 0.6328 - val



```
Epoch 100/100
276/276 [=====] - 28s 99ms/step - loss: 1.4391 - accuracy: 0.6336 - val_
```

11

```
12 # Compile the model
    from keras import metrics

    model.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy',
                          metrics.AUC(name='my_auc'),
                          F1_Score])

-----

NameError                                Traceback (most recent call last)

<ipython-input-12-da8c9a7433e9> in <module>
      6             metrics=['accuracy',
      7                 metrics.AUC(name='my_auc'),
----> 8                 F1_Score])
      9

NameError: name 'F1_Score' is not defined


# https://keras.io/api/callbacks/
# We can use a variety of pre-defined callbacks.
# Experiment with ReduceLROnPlateau()

import tensorflow_addons as tfa

from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau, CSVLogger, LearningRateScheduler

# We can also do a modelcheck point
# https://machinelearningmastery.com/check-point-deep-learning-models-keras/

# checkpoint to save the model with best validation accuracy
checkpoint = ModelCheckpoint(filepath='model.{epoch:02d}-{val_loss:.2f}.h5',
                             monitor='val_accuracy', verbose=1, save_best_only=True, mode='max')

# We can also stop the model early
#https://machinelearningmastery.com/how-to-stop-training-deep-neural-networks-at-the-right-time-when-the-training-is-over/
# val_loss
early_stopping_callback = EarlyStopping(monitor='accuracy', mode='min', verbose=1, patience=200)

# initialize TimeStopping callback
# https://www.tensorflow.org/addons/tutorials/time_stopping
# note that it will still run a minimum of 1 epoch
time_stopping_callback = tfa.callbacks.TimeStopping(seconds=600, verbose=1)

# We can also use CSVLogger to log information in a CSV
csvlogger = CSVLogger("logfile.csv", separator=',', append=False)

# ** IMPORTANT ** - please make sure that csvlogger is the last call back
# in the list.

my_callbacks = [time_stopping_callback, early_stopping_callback, checkpoint, csvlogger]
```

```

# Fitting the model with call-backs

num_epochs = 1

history=model.fit(train_data,
                  steps_per_epoch =train_steps_per_epoch,
                  validation_data=valid_data,
                  epochs=num_epochs,
                  validation_steps=valid_steps_per_epoch,
                  callbacks=my_callbacks)


# Compile the model
from keras import metrics
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy',
                      metrics.AUC(name='auc'),
                      metrics.Precision(name='precision'),
                      metrics.Recall(name='recall')])


# Fitting the model with more metrics

num_epochs = 1

history=model.fit(train_data,
                  steps_per_epoch =train_steps_per_epoch,
                  validation_data=valid_data,
                  epochs=num_epochs,
                  validation_steps=valid_steps_per_epoch,
                  callbacks=my_callbacks)


# Defining custom metrics to record while running
from keras import backend as K

def F1_Score(y_true, y_pred): #taken from old keras source code
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    recall = true_positives / (possible_positives + K.epsilon())
    f1_val = 2*(precision*recall)/(precision+recall+K.epsilon())
    return f1_val

def my_metric_fn(y_true, y_pred):
    squared_difference = tf.square(y_true - y_pred)
    return tf.reduce_mean(squared_difference, axis=-1) # Note the `axis=-1`


# Compile the model
from keras import metrics
model.compile(loss='categorical_crossentropy',
              optimizer='adam',

```

```

        metrics=['accuracy',
                 metrics.AUC(name='auc'),
                 metrics.Precision(name='precision'),
                 metrics.Recall(name='recall'),
                 F1_Score])

# Fitting the model with more metrics

num_epochs = 1

history=model.fit(train_data,
                  steps_per_epoch=train_steps_per_epoch,
                  validation_data=valid_data,
                  epochs=num_epochs,
                  validation_steps=valid_steps_per_epoch,
                  callbacks=my_callbacks)

# Defining custom call backs

# https://www.tensorflow.org/guide/keras/custom\_callback
# https://keras.io/guides/writing\_your\_own\_callbacks/

from keras.callbacks import Callback
import time

class TimingCallback(keras.callbacks.Callback):
    def __init__(self):
        super(TimingCallback, self).__init__()
    def on_batch_begin(self, epoch, logs=None):
        self.starttime=time.time()
    def on_batch_end(self, epoch, logs=None):
        logs['epoch_time'] = (time.time()-self.starttime)
        print('\nepoch_time(sec)=' ,logs['epoch_time'],'\n')

# create an instance of the timingcallback
timing_call = TimingCallback()

# We can also use other metrics
# https://keras.io/api/metrics/
class PrintBatchCallback(keras.callbacks.Callback):
    def on_train_batch_end(self, batch, logs=None):
        print("For batch {}, loss is {:.2f}.".format(batch, logs["loss"]))
        print("For batch {}, accuracy is {:.2f}.".format(batch, logs["accuracy"]))
        print("For batch {}, AUC is {:.2f}.".format(batch, logs["auc"]))

print_batch_call = PrintBatchCallback()

# add to the callback list
my_callbacks = [time_stopping_callback,early_stopping_callback,checkpoint,print_batch_call, timing_call]

# Compile the model
from keras import metrics
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy',
                      metrics.AUC(name='auc'),
                      metrics.Precision(name='precision'),
                      metrics.Recall(name='recall'),
                      F1_Score])

# Fitting the model with more metrics

```

```

num_epochs = 1

history=model.fit(train_data,
                  steps_per_epoch =train_steps_per_epoch,
                  validation_data=valid_data,
                  epochs=num_epochs,
                  validation_steps=valid_steps_per_epoch,
                  callbacks=my_callbacks)

# https://neptune.ai/blog/keras-metrics

# How to save batch level data in a file

import os
from keras.callbacks import Callback
import numpy as np

class SaveBatchLevelDataCallback(keras.callbacks.Callback):
    def __init__(self, validation_data, save_dir):
        super().__init__()
        self.validation_data = validation_data
        os.makedirs(save_dir, exist_ok=True)
        self.save_dir = save_dir
        self.f = None

    def on_epoch_begin(self, epoch, logs=None):
        # create a file
        self.f= open(os.path.join(self.save_dir, f'epoch_{epoch}.csv'),'w+')
        line = "batch,loss,accuracy,auc\n"
        self.f.write(line)

    def on_epoch_end(self, batch, logs=None):
        self.f.close()

    def on_train_batch_end(self, batch, logs=None):
        line = "{},{:7.2f},{:7.2f},{:7.2f}\n".format(batch, logs["loss"], logs["accuracy"],logs[
        self.f.write(line)

batch_write_cbk = SaveBatchLevelDataCallback(validation_data=valid_data,save_dir='batch_data')

# add to the callback list
my_callbacks = [time_stopping_callback,early_stopping_callback,checkpoint,batch_write_cbk, CSVLo

# # Compile the model
from keras import metrics
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy',
                      metrics.AUC(name='auc'),
                      metrics.Precision(name='precision'),
                      metrics.Recall(name='recall'),
                      F1_Score])

# Fitting the model with more metrics

num_epochs = 10

history=model.fit(train_data,

```

```
steps_per_epoch=train_steps_per_epoch,  
validation_data=valid_data,  
epochs=num_epochs,  
validation_steps=valid_steps_per_epoch,  
callbacks=my_callbacks)
```

```
# print history  
print(history.history)
```

```
#plot accuracy vs epoch  
plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.title('Model accuracy')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validate'], loc='upper left')  
plt.show()
```

```
# Plot loss values vs epoch  
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Model loss')  
plt.ylabel('Loss')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validate'], loc='upper left')  
plt.show()
```

```
# Plot loss values vs epoch  
plt.plot(history.history['F1_Score'])  
plt.plot(history.history['val_F1_Score'])  
plt.title('Model F1-Score')  
plt.ylabel('F1_Score')  
plt.xlabel('Epoch')  
plt.legend(['Train', 'Validate'], loc='upper left')  
plt.show()
```

```
# Plot accuracy vs. prevision  
plt.plot(history.history['precision'],label='precision')  
plt.plot(history.history['val_precision'],label='val_precision')  
plt.plot(history.history['recall'],label='recall')  
plt.plot(history.history['val_recall'],label='val_precision')  
plt.title('Model Precision and Recall')  
plt.ylabel('Precision and Recall')  
plt.xlabel('Epoch')  
plt.legend()  
plt.show()
```

```
# Plot accuracy vs. prevision  
plt.plot(history.history['precision'],history.history['recall'],'o', color='black',label='precis')  
plt.plot(history.history['recall'],history.history['val_recall'],'o', color='red',label='val_pre')  
plt.title('Model Precision and Recall')  
plt.ylabel('Precision')  
plt.xlabel('Recall')  
plt.legend()  
plt.show()
```

```
# Evaluate against test data.  
scores = model.evaluate(test_data, verbose=1)
```

```
print('Test loss:', scores[0])  
print('Test accuracy:', scores[1])  
print('Test AUC:', scores[1])
```

```

print('Test precision:', scores[1])
print('Test recall:', scores[1])
print('Test F1-Score:', scores[1])

# For evaluation first, we will create the actual and predicted labels
# We can then use these to generate all the reports we need.

# make predictions on the testing images, finding the index of the
# label with the corresponding largest predicted probability

predicted = model.predict(x=test_data, steps=test_steps_per_epoch)

# create predicted IDs
predicted = np.argmax(predicted, axis=1)

# create test labels from the generator
actual = []
for i in range(0,int(test_steps_per_epoch)):
    actual.extend(np.array(test_data[i][1]))

# create actual IDs
actual = np.asarray(actual).argmax(axis=1)

# make sure predicted and actual are the same size and shape
print(predicted.shape)
print(actual.shape)

from sklearn.metrics import classification_report

print("[INFO] evaluating network...")
print(classification_report(actual, predicted))

# Now we can determine the confusion matrix
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(actual,predicted)

def print_cm(cm, frm, to,abs_or_relative=0):
    import seaborn as sns
    import matplotlib.pyplot as plt

    cm = cm[frm:to+1,frm:to+1]
    # create labels
    x_axis_labels = np.arange(frm,to+1)
    y_axis_labels = np.arange(frm,to+1)

    plt.xticks(rotation=45)
    plt.yticks(rotation=-45)

    if(abs_or_relative==0):
        sns.heatmap(cm, annot=True,xticklabels=x_axis_labels, yticklabels=y_axis_labels)
    else:
        sns.heatmap(cm/np.sum(cm), annot=True,
            fmt='.2%', cmap='Blues',
            xticklabels=x_axis_labels, yticklabels=y_axis_labels)

print_cm(cm,1 ,20,0)

# we already have actual and predicted

```

```

# also see https://www.dlology.com/blog/simple-guide-on-how-to-generate-roc-plot-for-keras-class
# for micro-average ROC curves as well

import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc

fpr = dict()
tpr = dict()
roc_auc = dict()

#extract the actual labels from the test data
Y_test = []
for i in range(0,int(test_steps_per_epoch)):
    Y_test.extend(np.array(test_data[i][1]))
Y_test = np.array(Y_test)
n_classes = Y_test.shape[1] # one hot encoded

# create actual output from the model using test_data
y_score=model.predict(x=test_data, steps=test_steps_per_epoch)

print(Y_test.shape)
print(y_score.shape)

print(n_classes)
# compare each class's probabilities one by one
# each acts like a single column
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(Y_test[:,i], y_score[:,i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Print the AUC scores
from IPython.display import display
import pandas as pd
auc_array = np.array(list(roc_auc.items()))
df = pd.DataFrame(auc_array[:,1])
df.columns = ['AUC']
display(df)

# plot the ROC for the ith class cls
import matplotlib.pyplot as plt
import os

def plot_roc(cls,roc_dir):
    plt.plot(fpr[cls], tpr[cls], lw=2,label='ROC curve of class {0} (area = {1:0.3f})'
            ''.format(cls, roc_auc[cls]))
    plt.plot([0, 1], [0, 1], 'k--', lw=2)
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC')
    plt.legend(loc="lower right")
    plt.tight_layout()
    plt.savefig(os.path.join(roc_dir, f'ROC_{cls}.png'))
    plt.show()

# make sure directory exists
def make_directory(roc_dir):
    try:
        os.mkdir(roc_dir)

```

```

except OSError:
    print ("Creation of the directory %s failed" % roc_dir)
else:
    print ("Successfully created the directory %s " % roc_dir)

# print the roc curve for 0

make_directory('rocs')

for i in range(n_classes):
    plot_roc(i, 'rocs')

# Using tensorflow extension
# Load the TensorBoard notebook extension
%load_ext tensorboard
import datetime

# Define tensorboard callback

log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1)

# Using remote tensorboard
#https://blog.yyliu.net/remote-tensorboard/

# Compile the model
from keras import metrics
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy',
                      metrics.AUC(name='auc'),
                      metrics.Precision(name='precision'),
                      metrics.Recall(name='recall')])

# Fitting the model with more metrics

num_epochs = 10

history=model.fit(train_data,
                  steps_per_epoch=train_steps_per_epoch,
                  validation_data=valid_data,
                  epochs=num_epochs,
                  validation_steps=valid_steps_per_epoch,
                  callbacks=[tensorboard_callback])

#%tensorboard --logdir logs/fit

```



